

DAT2 week 9


Misrab

Agenda

- Learning Theory
- Reinforcement Learning
- CS
- HMMs

Learning Theory

Bias / Variance: parameters, models



How much data needed?

Which models to even consider?

Learning Theory

Lemma. (The union bound). Let A_1, A_2, \dots, A_k be k different events (that may not be independent). Then

$$P(A_1 \cup \dots \cup A_k) \leq P(A_1) + \dots + P(A_k).$$

Lemma. (Hoeffding inequality) Let Z_1, \dots, Z_m be m independent and identically distributed (iid) random variables drawn from a Bernoulli(ϕ) distribution. I.e., $P(Z_i = 1) = \phi$, and $P(Z_i = 0) = 1 - \phi$. Let $\hat{\phi} = (1/m) \sum_{i=1}^m Z_i$ be the mean of these random variables, and let any $\gamma > 0$ be fixed. Then

$$P(|\phi - \hat{\phi}| > \gamma) \leq 2 \exp(-2\gamma^2 m)$$

How many training points (samples) to be within 0.01 of each other with probability 95%?

Learning Theory

Binary classification empirical error

$$\hat{\varepsilon}(h) = \frac{1}{m} \sum_{i=1}^m 1\{h(x^{(i)}) \neq y^{(i)}\}.$$

(fraction misclassified)

Generalisation error

$$\varepsilon(h) = P_{(x,y) \sim \mathcal{D}}(h(x) \neq y).$$

Learning Theory

PAC framework
“probably approximately correct”

i.e. training and testing from same distribution

Learning Theory

Minimising training error

$$\hat{\theta} = \arg \min_{\theta} \hat{e}(h_{\theta}).$$

$$\mathcal{H} = \{h_{\theta} : h_{\theta}(x) = 1\{\theta^T x \geq 0\}, \theta \in \mathbb{R}^{n+1}\}$$

equivalent to $\hat{h} = \arg \min_{h \in \mathcal{H}} \hat{e}(h)$

Learning Theory

Finite $\mathcal{H} = \{h_1, \dots, h_k\}$

We want guarantees on generalisation error when we minimise empirical error

- theoretical framework
- applied cases e.g. too few data points for cross-validation

Learning Theory

- Take fixed h_i in H
- Let $Z_j = 1\{h_i(x^{(j)}) \neq y^{(j)}\}$.
- Training error of h_i

$$\hat{\varepsilon}(h_i) = \frac{1}{m} \sum_{j=1}^m Z_j.$$

- Hoeffding

$$P(|\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) \leq 2 \exp(-2\gamma^2 m).$$

Learning Theory

Great. But we don't just want this for one h_i , we want it for all of H simultaneously

$$P(A_i) \dots \blacktriangledown P(|\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) \leq 2 \exp(-2\gamma^2 m).$$

$$\begin{aligned} P(\exists h \in \mathcal{H}. |\varepsilon(h) - \hat{\varepsilon}(h)| > \gamma) &= P(A_1 \cup \dots \cup A_k) \\ &\leq \sum_{i=1}^k P(A_i) \\ &\leq \sum_{i=1}^k 2 \exp(-2\gamma^2 m) \\ &= 2k \exp(-2\gamma^2 m) \end{aligned}$$

Learning Theory

$$\begin{aligned}P(\exists h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) &= P(A_1 \cup \dots \cup A_k) \\&\leq \sum_{i=1}^k P(A_i) \\&\leq \sum_{i=1}^k 2 \exp(-2\gamma^2 m) \\&= 2k \exp(-2\gamma^2 m)\end{aligned}$$

If we subtract both sides from 1, we find that

$$\begin{aligned}P(\neg \exists h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| > \gamma) &= P(\forall h \in \mathcal{H}. |\varepsilon(h_i) - \hat{\varepsilon}(h_i)| \leq \gamma) \\&\geq 1 - 2k \exp(-2\gamma^2 m)\end{aligned}$$

Learning Theory

Example

For instance, we can ask the following question: Given γ and some $\delta > 0$, how large must m be before we can guarantee that with probability at least $1 - \delta$, training error will be within γ of generalization error? By setting $\delta = 2k \exp(-2\gamma^2 m)$ and solving for m , [you should convince yourself this is the right thing to do!], we find that if

$$m \geq \frac{1}{2\gamma^2} \log \frac{2k}{\delta},$$

then with probability at least $1 - \delta$, we have that $|\varepsilon(h) - \hat{\varepsilon}(h)| \leq \gamma$ for all $h \in \mathcal{H}$. (Equivalently, this shows that the probability that $|\varepsilon(h) - \hat{\varepsilon}(h)| > \gamma$ for some $h \in \mathcal{H}$ is at most δ .) This bound tells us how many training examples we need in order make a guarantee. The training set size m that a certain method or algorithm requires in order to achieve a certain level of performance is also called the algorithm's **sample complexity**.

Learning Theory

With just a bit more effort, we can get an explicit guarantee on generalisation error

Theorem. Let $|\mathcal{H}| = k$, and let any m, δ be fixed. Then with probability at least $1 - \delta$, we have that

$$\varepsilon(\hat{h}) \leq \left(\min_{h \in \mathcal{H}} \varepsilon(h) \right) + 2\sqrt{\frac{1}{2m} \log \frac{2k}{\delta}}.$$

lower bias: increasing
H class

variance tradeoff of
increasing H class

Learning Theory

In summary, given a hypothesis class of complexity k , we can figure out how many training examples we need to be within a certain distance from minimising generalisation error when carrying out empirical risk minimisation

Analog for continuous case:
Vapnik-Chervonenki (VC) dimension

Reinforcement Learning

Previously, y had a “correct” value

Now instead, what if the best we can provide is a reward function

e.g. robot motion, gaming...trading?

<https://www.youtube.com/watch?v=qv6UVQQ0F44>

<https://www.youtube.com/watch?v=M8YjvHYbZ9w>

<https://www.youtube.com/watch?v=rVlhMGQgDkY>

Reinforcement Learning

A Markov decision process is a tuple $(S, A, \{P_{sa}\}, \gamma, R)$, where:

- S is a set of **states**. (For example, in autonomous helicopter flight, S might be the set of all possible positions and orientations of the helicopter.)
- A is a set of **actions**. (For example, the set of all possible directions in which you can push the helicopter's control sticks.)
- P_{sa} are the state transition probabilities. For each state $s \in S$ and action $a \in A$, P_{sa} is a distribution over the state space. We'll say more about this later, but briefly, P_{sa} gives the distribution over what states we will transition to if we take action a in state s .
- $\gamma \in [0, 1)$ is called the **discount factor**.
- $R : S \times A \mapsto \mathbb{R}$ is the **reward function**. (Rewards are sometimes also written as a function of a state S only, in which case we would have $R : S \mapsto \mathbb{R}$).

Reinforcement Learning

$$s_0 \xrightarrow{a_0} s_1 \xrightarrow{a_1} s_2 \xrightarrow{a_2} s_3 \xrightarrow{a_3} \dots$$
$$s_1 \sim P_{s_0 a_0}.$$

Upon visiting the sequence of states s_0, s_1, \dots with actions a_0, a_1, \dots , our total payoff is given by

$$R(s_0, a_0) + \gamma R(s_1, a_1) + \gamma^2 R(s_2, a_2) + \dots .$$

Or, when we are writing rewards as a function of the states only, this becomes

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots .$$

Reinforcement Learning

We want to maximise expected discounted reward

Economic interpretation of gamma?

$$E [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots]$$

Reinforcement Learning

Policy to choose actions

$$\pi : S \mapsto A$$

$$a = \pi(s)$$

Value function based on policy

$$V^\pi(s) = \mathbb{E} [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots \mid s_0 = s, \pi].$$

Reinforcement Learning

$$V^\pi(s) = \mathbb{E} [R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots \mid s_0 = s, \pi].$$

Bellman equation

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s').$$

Can you draw this (perhaps as a decision tree?)

Reinforcement Learning

We want to find the optimal value

$$V^*(s) = \max_{\pi} V^{\pi}(s).$$

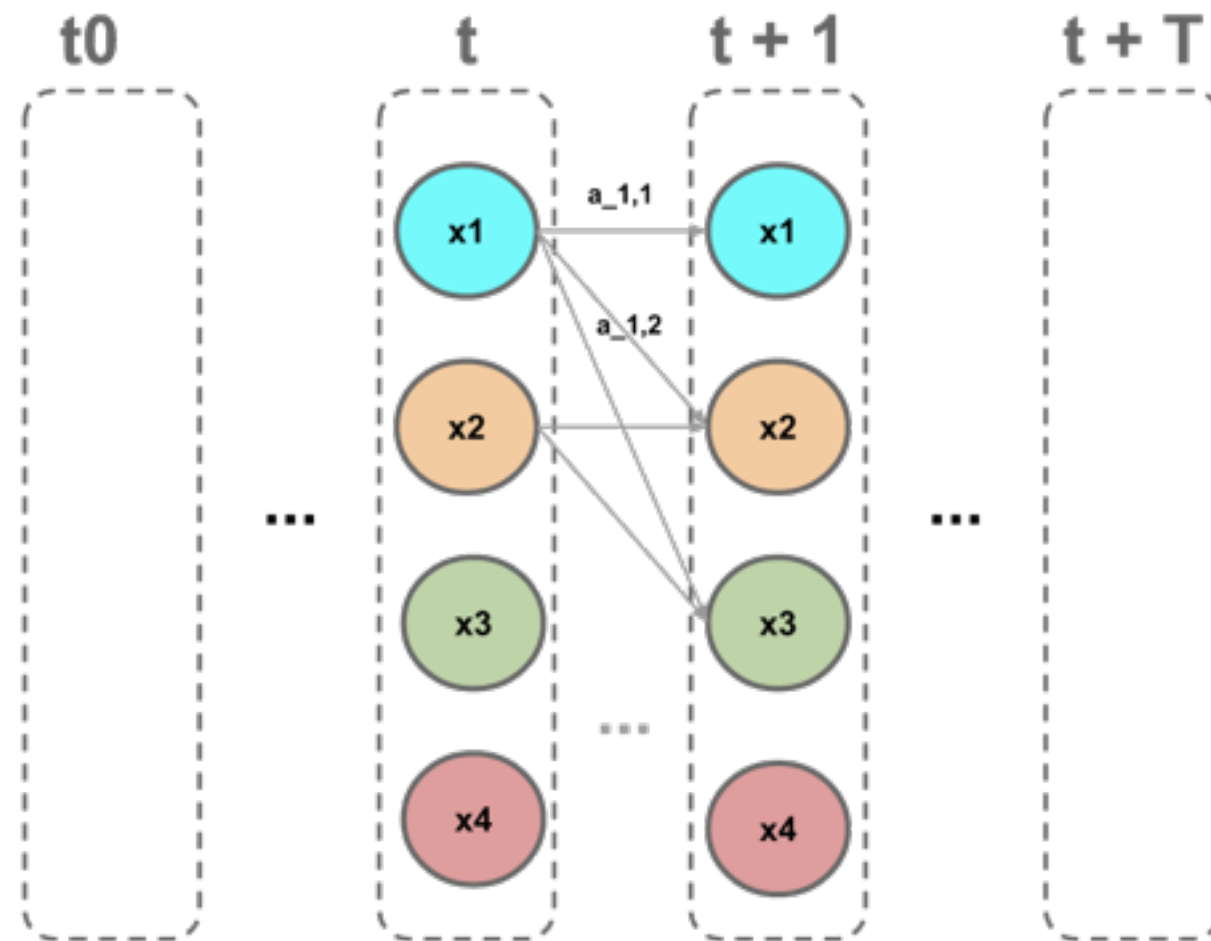
or

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s').$$

max over possible actions

Reinforcement Learning

Exponential paths to explore



Hansel and Grettel!

Reinforcement Learning

Recommended for Dynamic Programming and
Value Iteration:

<https://en.wikipedia.org/wiki/Memoization>

<https://www.youtube.com/watch?v=oefOCk3koZo>

<https://www.youtube.com/watch?v=ip4iSMRW5X4>

Reinforcement Learning

Extensions:

Learning P_{sa}

Continuous space (discretisation)

Computer Science

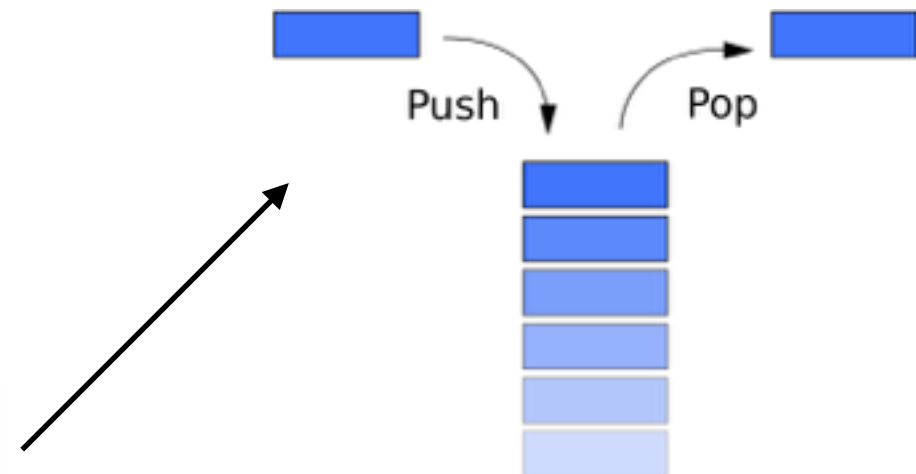
Scope and the call stack

```
def f():  
    x = 2  
    return x + 4
```

```
x = 3  
print f()
```

what is printed?

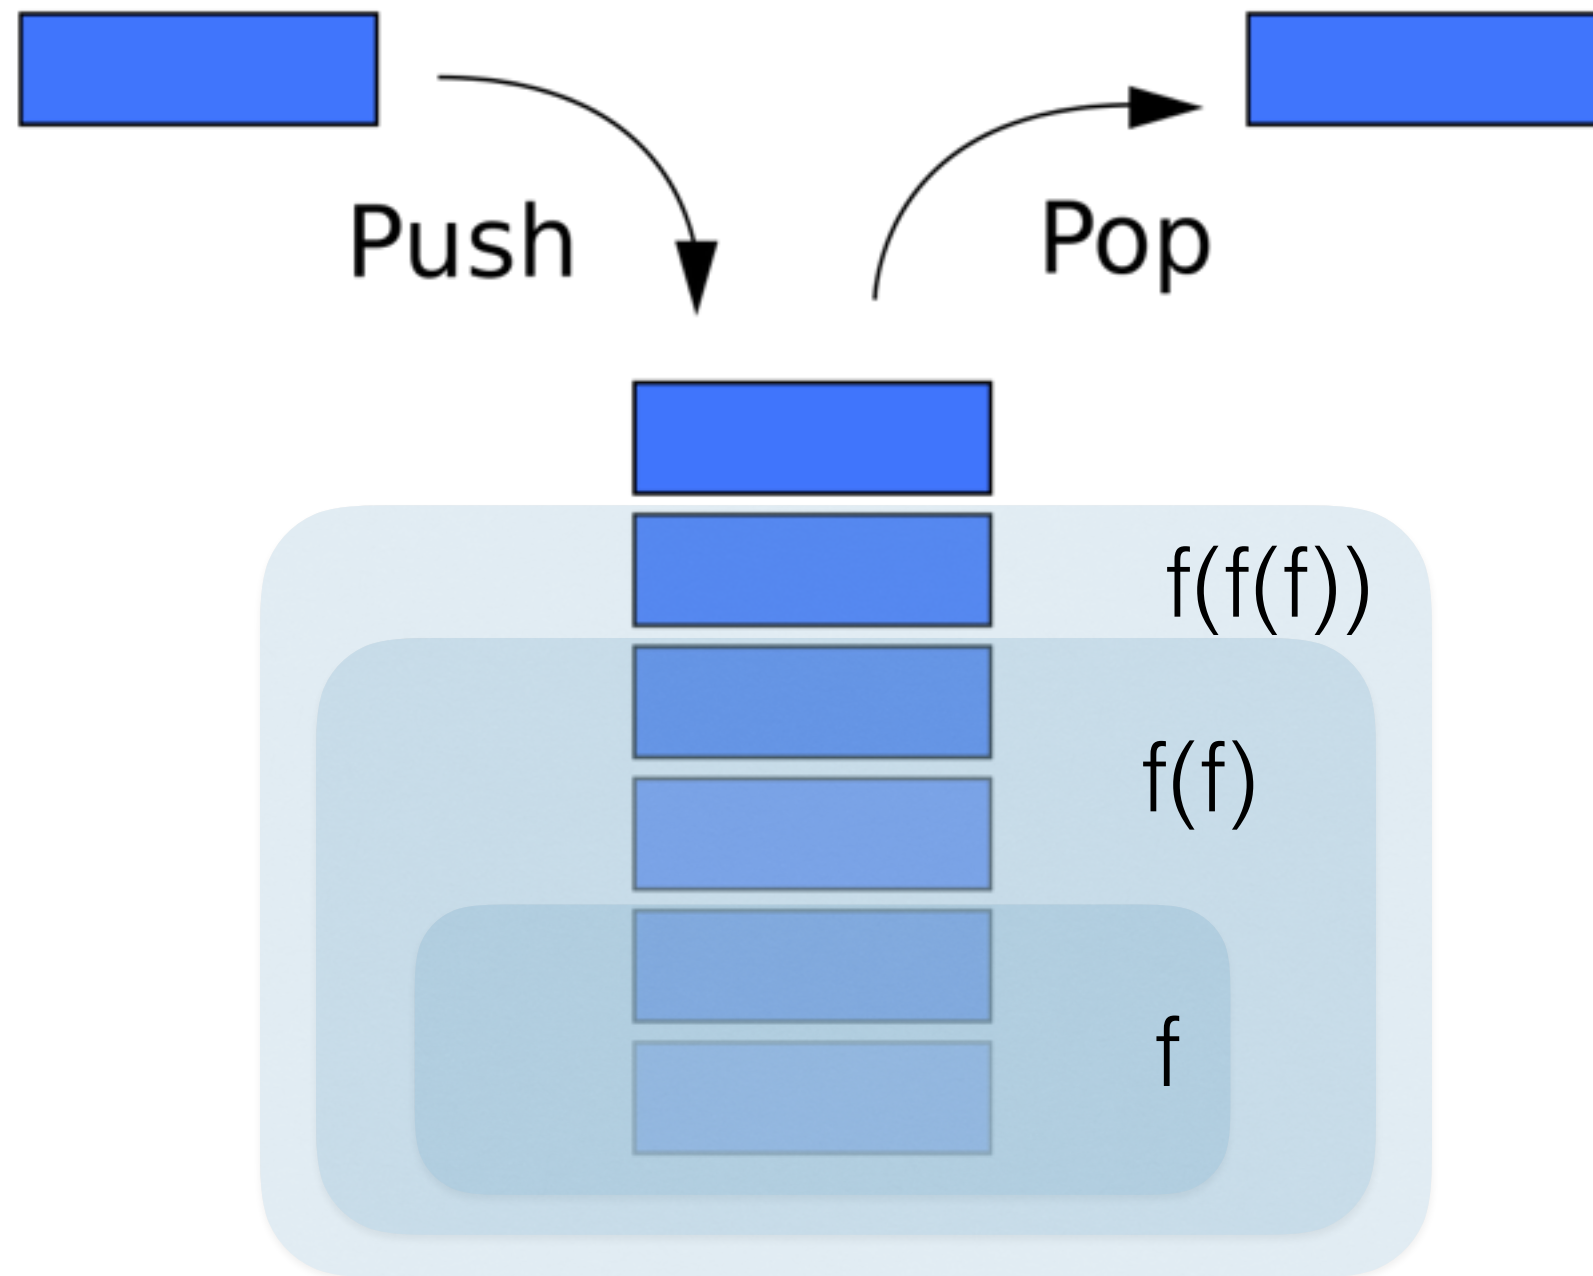
```
def f():  
    x = 2  
    return x + 4
```




```
x = 3  
print f()
```

what is printed?

Recursion: function calling itself



Pointers!



How pythons can bite!

```
arr = [ { ... } , { ... } , { ... } ]  
results = []
```

```
for ...  
    x = arr[i]  
    x.moo = "foo"
```

```
results.append(x)
```

What's wrong?

...

results = [x, x, x, x]
x —> the last item

```
arr = [ { ... } , { ... } , { ... } ]  
results = []
```

```
for ...
```

```
    x = copy.deepcopy(arr[i])
```

```
    x.moo = "foo"
```

```
    results.append(x)
```

Practice!

```
int *moo // pointer  
int moo // actual value
```

```
int moo = 3  
&moo // address of value
```

```
int *moo  
*moo = 3 // set value!
```

http://www.gdsw.at/languages/c/programming-bbrownc_0771.htm

Hidden Markov Models

Sequence of words over time based on audio

Other examples?

Easy case: we observe the actual state

$$S = \{ \text{rain, sun, cloud} \}$$

Observed over time:

$$z_1 = \text{sun}, z_2 = \text{cloud}, z_3 = \text{sun} \dots$$

$$P(z_t | z_{t-1}, z_{t-2}, \dots, z_1) = P(z_t | z_{t-1})$$

Simplifying Assumptions

Limited horizon: $P(z_t|z_{t-1}, z_{t-2}, \dots, z_1) = P(z_t|z_{t-1})$

Stationary process: $P(z_t|z_{t-1}) = P(z_2|z_1); t \in 2...T$

Can you explain this in a picture?

State transition matrix

$$A = \begin{array}{cc} & \begin{array}{c} s_0 \quad s_{sun} \quad s_{cloud} \quad s_{rain} \end{array} \\ \begin{array}{c} s_0 \\ s_{sun} \\ s_{cloud} \\ s_{rain} \end{array} & \begin{array}{ccccc} 0 & .33 & .33 & .33 \\ 0 & .8 & .1 & .1 \\ 0 & .2 & .6 & .2 \\ 0 & .1 & .2 & .7 \end{array} \end{array}$$

What's the prob of a given sequence?

$$\begin{aligned}P(\vec{z}) &= P(z_t, z_{t-1}, \dots, z_1; A) \\&= P(z_t, z_{t-1}, \dots, z_1, z_0; A) \\&= P(z_t|z_{t-1}, z_{t-2}, \dots, z_1; A)P(z_{t-1}|z_{t-2}, \dots, z_1; A) \dots P(z_1|z_0; A) \\&= P(z_t|z_{t-1}; A)P(z_{t-1}|z_{t-2}; A) \dots P(z_2|z_1; A)P(z_1|z_0; A)\end{aligned}$$

$$\begin{aligned}&= \prod_{t=1}^T P(z_t|z_{t-1}; A) \\&= \prod_{t=1}^T A_{z_{t-1} z_t}\end{aligned}$$

Compute this:

$$P(z_1 = s_{sun}, z_2 = s_{cloud}, z_3 = s_{rain}, z_4 = s_{rain}, z_5 = s_{cloud})$$


$$A = \begin{array}{cc} & \begin{array}{c} s_0 \\ s_{sun} \\ s_{cloud} \\ s_{rain} \end{array} \\ \begin{array}{c} s_0 \\ s_{sun} \\ s_{cloud} \\ s_{rain} \end{array} & \begin{array}{ccccc} s_0 & s_{sun} & s_{cloud} & s_{rain} \\ 0 & .33 & .33 & .33 \\ 0 & .8 & .1 & .1 \\ 0 & .2 & .6 & .2 \\ 0 & .1 & .2 & .7 \end{array} \end{array}$$

If we had a sequence **z**, we could use maximum likelihood to figure out **A**

Examples?

Hidden Markov Models

We don't observe the sequence directly



e.g. words vs audio waves

New matrix **B** that also tells us $P(y_t = i \mid x_t = j)$

use observe sequence **z**

$$\begin{aligned} P(\vec{x}; A, B) &= \sum_{\vec{z}} P(\vec{x}, \vec{z}; A, B) \\ &= \sum_{\vec{z}} P(\vec{x}|\vec{z}; A, B) P(\vec{z}; A, B) \end{aligned}$$

use HMM assumptions

$$\begin{aligned} P(\vec{x}; A, B) &= \sum_{\vec{z}} P(\vec{x}|\vec{z}; A, B) P(\vec{z}; A, B) \\ &= \sum_{\vec{z}} \left(\prod_{t=1}^T P(x_t|z_t; B) \right) \left(\prod_{t=1}^T P(z_t|z_{t-1}; A) \right) \\ &= \sum_{\vec{z}} \left(\prod_{t=1}^T B_{z_t x_t} \right) \left(\prod_{t=1}^T A_{z_{t-1} z_t} \right) \end{aligned}$$

This is relatively advanced. Even sklearn
has outsourced it:

<https://github.com/hmmlearn/hmmlearn>